

# \* Marketing Camping \*

## This Is the Project Source Code

### Source Code :-:

```
#!/usr/bin/env python
# -*- coding: utf-8 -*-

# # Project: Marketing Campaigns – EDA & Hypothesis
# Testing

# # Step 1: imports

# Cell 2 - imports
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import cross_val_score
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import StandardScaler

pd.options.display.max_columns = 200
```

```
# ### Load data & quick checks

# Cell 3 - load and quick checks (adjust path if needed)
df = pd.read_csv('marketing_data.csv')      # replace with
correct path if needed

print("Rows, cols:", df.shape)
display(df.head())
display(df.info())

# ### Find & normalize key column names

# Cell 4 - helper to find columns by substring

def find_cols(df, substr):
    return [c for c in df.columns if substr.lower() in
c.lower()]

def find_first(df, *subs):
    for s in subs:
        cols = find_cols(df, s)
        if cols:
            return cols[0]
    return None

dt_col = find_first(df, 'dt_customer', 'date')
income_col = find_first(df, 'income')
edu_col = find_first(df, 'educ', 'education')
mar_col = find_first(df, 'marital', 'mar')
year_col = find_first(df, 'year_birth', 'birth', 'year')
kid_col = find_first(df, 'kid')
teen_col = find_first(df, 'teen')

print("Detected cols:", dt_col, income_col, edu_col, mar_col,
year_col, kid_col, teen_col)
```

```
# ### Parse dates & clean Income

# Cell 5 - parse dates and clean income

if dt_col:
    df[dt_col] = pd.to_datetime(df[dt_col], errors='coerce')

if income_col:
    df[income_col] = (df[income_col].astype(str)
                      .str.replace(',', '', regex=False)
                      .str.replace(' ', '', regex=False)
                      .str.replace(r'^[^\d.-]', '', regex=True))

    df[income_col] = pd.to_numeric(df[income_col],
                                   errors='coerce')

print("Income missing before imputation:",
      df[income_col].isna().sum() if income_col else 'Income not found')

df[[income_col]].describe()
```

```
# ### Clean categories (education & marital)

# Cell 6 - clean categorical columns

def clean_cat(s):
    return
s.astype(str).str.strip().str.title().replace({'Nan':'Unknown',
,'N/A':'Unknown'})

if edu_col:
    print("Education - before:", df[edu_col].unique()[:20])
    df[edu_col] = clean_cat(df[edu_col])
    print("Education - after:", df[edu_col].unique()[:20])

if mar_col:
    print("Marital - before:", df[mar_col].unique()[:20])
    df[mar_col] = clean_cat(df[mar_col])
    print("Marital - after:", df[mar_col].unique()[:20])
```

```
# ### Income imputation by (Education x Marital)

# Cell 7 - groupwise mean imputation (education x marital)
if income_col and edu_col and mar_col:
    grp_mean = df.groupby([edu_col,
    mar_col])[income_col].transform('mean')
    before = df[income_col].isna().sum()
    df[income_col] = df[income_col].fillna(grp_mean)
    df[income_col] =
df[income_col].fillna(df[income_col].median()) # fallback
    after = df[income_col].isna().sum()
    print(f"Income missing before: {before}; after imputation:
{after}")
    display(df[[edu_col, mar_col, income_col]].head())
else:
    print("Required columns for groupwise imputation not all
found.")
```

```

# ### Feature engineering (age, total_children,
total_spending, total_purchases)

# Cell 8 - feature engineering

# total_children

if kid_col or teen_col:
    k = kid_col if kid_col else None
    t = teen_col if teen_col else None
    df['total_children'] = 0
    if k:
        df['total_children'] += df[k].fillna(0).astype(float)
    if t:
        df['total_children'] += df[t].fillna(0).astype(float)
else:
    print("Kid/Teen columns not found - total_children remains 0")

# age (using year of birth if present)

if year_col and dt_col:
    try:
        if np.issubdtype(df[year_col].dtype, np.number):
            df['age'] = df[dt_col].dt.year - df[year_col]
        else:
            df[year_col] = pd.to_datetime(df[year_col],
                                         errors='coerce')
            df['age'] = ((df[dt_col] - df[year_col]).dt.days /
                         365.25).astype(int)
    except Exception as e:
        print("Age calc issue:", e)

# total_spending: sum columns with 'Mnt' or 'Amount' or
'Spend'

```

```
mnt_cols = [c for c in df.columns if any(k in c.lower() for k in ['mnt', 'amount', 'spend'])]

if not mnt_cols:
    # try product names often present
    mnt_cols = [c for c in df.columns if c.lower().startswith('mnt')]

df['total_spending'] = df[mnt_cols].sum(axis=1) if mnt_cols
else 0

# total_purchases: sum columns with 'Num' and 'Purch'
purchase_cols = [c for c in df.columns if 'num' in c.lower()
and 'purch' in c.lower()]

if not purchase_cols:
    # try common channel names
    purchase_cols = [c for c in df.columns if any(s in
c.lower() for s in ['web', 'catalog', 'store', 'purch',
'purchase']) and df[c].dtype != object]

df['total_purchases'] = df[purchase_cols].sum(axis=1) if
purchase_cols else 0

print("Features created. Sample:")
display(df.head()[['age', 'total_children', 'total_spending', 'total_purchases']].head())
print("Spending cols used:", mnt_cols)
print("Purchase cols used:", purchase_cols)
```

```
# ### EDA: distributions, boxplots, correlation

# Cell 9 - EDA plots
plt.figure(figsize=(10,4))

if income_col:
    plt.subplot(1,2,1)
    sns.histplot(df[income_col].dropna(), kde=False, bins=30)
    plt.title('Income distribution')

if 'total_spending' in df.columns:
    plt.subplot(1,2,2)

    sns.histplot(df['total_spending'].replace(0,np.nan).dropna(),
    kde=False, bins=30)
    plt.title('Total spending distribution')

plt.tight_layout()
plt.show()

# Boxplot for total_spending
plt.figure(figsize=(8,3))
sns.boxplot(x=df['total_spending'])
plt.title('Total spending (boxplot)')
plt.show()

# Correlation heatmap (numeric)
num = df.select_dtypes(include=[np.number])
plt.figure(figsize=(10,8))
sns.heatmap(num.corr(), annot=False, cmap='coolwarm')
plt.title('Numeric correlation heatmap')
plt.show()
```

```
# ### Outlier treatment (IQR capping)

# Cell 10 - IQR capping for total_spending

def cap_iqr(series, k=1.5):
    q1 = series.quantile(0.25)
    q3 = series.quantile(0.75)
    iqr = q3 - q1
    low = q1 - k * iqr
    high = q3 + k * iqr
    return series.clip(lower=low, upper=high)

if 'total_spending' in df.columns:
    df['total_spending_capped'] =
cap_iqr(df['total_spending'].fillna(0))

    # show before/after
    plt.figure(figsize=(10,3))
    plt.subplot(1,2,1)
    sns.boxplot(x=df['total_spending'])
    plt.title('Before capping')
    plt.subplot(1,2,2)
    sns.boxplot(x=df['total_spending_capped'])
    plt.title('After capping')
    plt.tight_layout()
    plt.show()

else:
    print("No total_spending column found.")
```

```

# ### Hypothesis testing (t-test + chi-square)

# Cell 11 - t-test: Do high-income customers spend more?

if income_col and 'total_spending' in df.columns:
    median_inc = df[income_col].median()

    high = df[df[income_col] >
median_inc]['total_spending_capped'].dropna()

    low = df[df[income_col] <=
median_inc]['total_spending_capped'].dropna()

    tstat, pval = stats.ttest_ind(high, low, equal_var=False)
    print("Median income:", median_inc)

    print("High group size:", len(high), "Low group size:",
len(low))

    print("t-statistic:", tstat, "p-value:", pval)

    if pval < 0.05:
        print("Interpretation: p < 0.05 -> significant
difference in spending between income groups.")

    else:
        print("Interpretation: no significant difference at
α=0.05.")

else:
    print("Income or spending data not available for t-test.")

# Cell 12 - chi-square: response vs education

resp_col = find_first(df, 'response', 'resp')
if resp_col and edu_col:
    ct = pd.crosstab(df[resp_col].fillna('Unknown'),
df[edu_col].fillna('Unknown'))

    from scipy.stats import chi2_contingency
    chi2, p, dof, expected = chi2_contingency(ct)
    print("Chi2 p-value:", p)

    if p < 0.05:

```

```
    print("Interpretation: p < 0.05 -> response depends on  
education (reject independence).")  
  
else:  
  
    print("Interpretation: no evidence of dependence at  
 $\alpha=0.05.$ ")  
  
    display(ct)  
  
else:  
  
    print("Missing response or education column; cannot run  
chi-square.")
```

```

# ### Simple predictive model (logistic regression)

# Cell 13 - simple logistic regression for response
resp_col = find_first(df, 'response', 'resp')
if resp_col:
    # build feature list from engineered features and reliable
    numeric columns
    candidate_feats =
    ['Age','age','Income','income','total_spending_capped','total_
    spending','total_purchases','total_children']
    features = [f for f in candidate_feats if f in df.columns]
    if not features:
        # fallback to numeric columns
        features =
        df.select_dtypes(include=[np.number]).columns.tolist()
        features = [f for f in features if f != resp_col][:6]
    print("Features used:", features)
    X = df[features].fillna(0)
    # ensure binary response
    y = df[resp_col]
    # convert to binary if not numeric
    if y.dtype == object:
        y = (y != 'No') & (y != '0') & (y != '')
        y = y.astype(int)
    else:
        y = (y != 0).astype(int)
    pipe = make_pipeline(StandardScaler(),
LogisticRegression(max_iter=1000))
    scores = cross_val_score(pipe, X, y, cv=5,
scoring='accuracy')
    print("CV accuracy (5-fold):", scores.mean(), "±",
scores.std())
    # fit and show coefficients

```

```
pipe.fit(X, y)
coefs = pipe.named_steps['logisticregression'].coef_[0]
print("Feature coefficients (log-odds):")
for feat, c in zip(features, coefs):
    print(f"{feat}: {c:.4f}")
else:
    print("Response column not found - cannot run logistic
regression.")
```